

5 paired differences per task, linear probe: each grey line is one pretraining run
positive = worse than the 162-class model; bars are the 95% paired interval over seeds, from the pre-specified analysis

